

simmer

Discrete-Event Simulation for R



Iñaki Úcar, postdoctoral fellow @ UC3M-Santander Big Data Institute (IBiDat)

Bart Smeets, data scientist @ dataroots.io

November 15-16, 2019



Simulation

From R. Shannon (1975), simulation is

the process of designing a **model of a real system** and conducting experiments with this model for the purpose either of **understanding the behavior of the system** or of **evaluating various strategies** [...] for the operation of the system.



Simulation

From R. Shannon (1975), simulation is

the process of designing a **model of a real system** and conducting experiments with this model for the purpose either of **understanding the behavior of the system** or of **evaluating various strategies** [...] for the operation of the system.

Taxonomy, from Law and Kelton (2000):

1. deterministic vs. stochastic
2. (time component?) static vs. dynamic
3. (if dynamic) continuous vs. discrete



Simulation

From R. Shannon (1975), simulation is

the process of designing a **model of a real system** and conducting experiments with this model for the purpose either of **understanding the behavior of the system** or of **evaluating various strategies** [...] for the operation of the system.

Taxonomy, from Law and Kelton (2000):

1. deterministic vs. stochastic
2. (time component?) static vs. dynamic
3. (if dynamic) continuous vs. discrete

Examples:

- deterministic + dynamic + continuous = *Dynamical Systems*
- stochastic + static = *Monte Carlo Simulation*
- **stochastic + dynamic + discrete = *Discrete-Event Simulation (DES)***



Simulation

What can be modelled as a Discrete-Event Simulation (DES)?

DES

- An **event** is an instantaneous occurrence that may change the **state of the system**
- The number of events is countable
- Between events, all the state variables remain constant
- Output: snapshots of the state of the system over time



Simulation

DES

What can be modelled as a Discrete-Event Simulation (DES)?

- An **event** is an instantaneous occurrence that may change the **state of the system**
- The number of events is countable
- Between events, all the state variables remain constant
- Output: snapshots of the state of the system over time

Common examples:

- customers arriving at a bank,
- products being manipulated in a supply chain,
- packets traversing a network,
- ...

and many more applications from manufacturing systems, construction engineering, project management, logistics, transportation systems, business processes, healthcare, telecommunications networks...



Simulation

Programming styles (Banks 2005):

DES

- **Activity-oriented**: fixed time increments; scan and verify conditions
- **Event-oriented**: event routines; event list ordered by time of occurrence
- **Process-oriented**: life cycle of entities of the real system; triggered by events



Simulation

DES

Programming styles (Banks 2005):

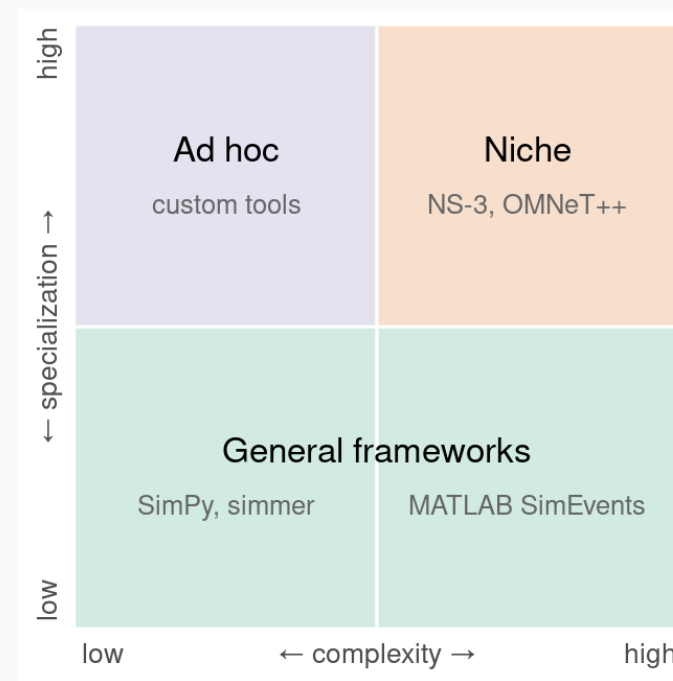
- **Activity-oriented**: fixed time increments; scan and verify conditions
- **Event-oriented**: event routines; event list ordered by time of occurrence
- **Process-oriented**: life cycle of entities of the real system; triggered by events

Spectrum of tools:

- High **complexity** and **specialization** generally means more accuracy

but

- More specialization requires more effort
- More complexity requires more effort





Simulation

DES

`simmer`

Main characteristics:

- General versatile framework for fast prototyping
- Rich and user-friendly R API over a fast C++ simulation core
- Process-oriented **trajectory-based** modelling
- Automatic monitoring capabilities
- Integration with R: repeatability, analysis, visualization





Simulation

DES

simmer

Main characteristics:

- General versatile framework for fast prototyping
- Rich and user-friendly R API over a fast C++ simulation core
- Process-oriented **trajectory-based** modelling
- Automatic monitoring capabilities
- Integration with R: repeatability, analysis, visualization

Resources:

- Online documentation (manual + 10 vignettes): <https://r-simmer.org>
- **Ucar I**, Smeets B, Azcorra A (2019). “simmer: Discrete-Event Simulation for R.” *Journal of Statistical Software*, 90(2), 1-30. doi: [10.18637/jss.v090.i02](https://doi.org/10.18637/jss.v090.i02).
- **Ucar I**, Hernández JA, Serrano P, Azcorra A (2018). “Design and Analysis of 5G Scenarios with simmer: An R Package for Fast DES Prototyping.” *IEEE Communications Magazine*, 56(11), 145-151. doi: [10.1109/MCOM.2018.1700960](https://doi.org/10.1109/MCOM.2018.1700960).

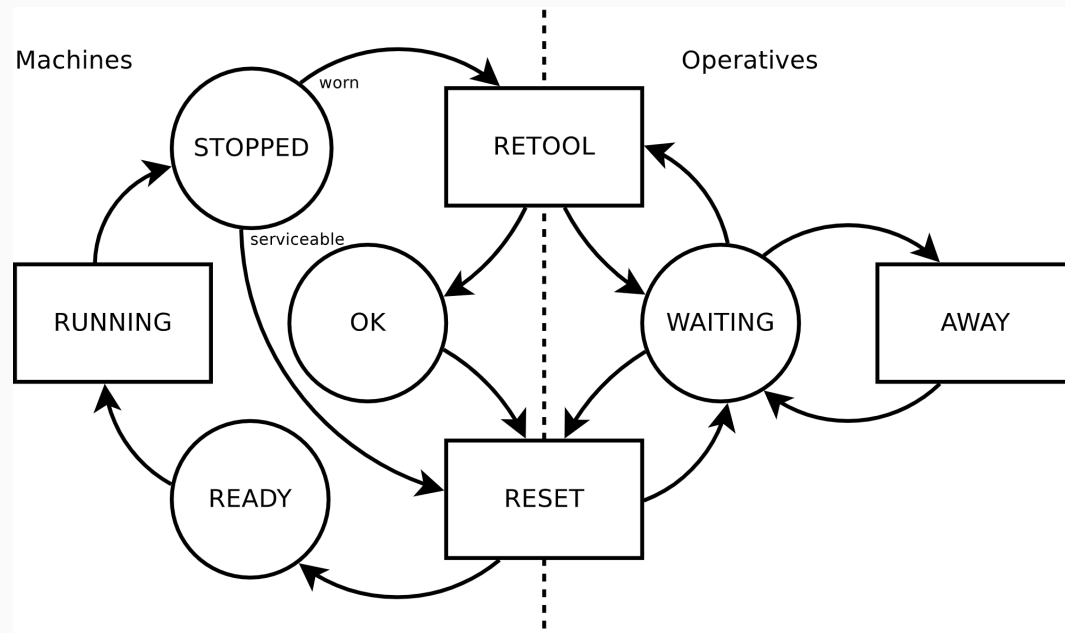


Example: Job Shop



From M. Pidd (1988), Section 5.3.1:

- Jobs are allocated to the first available machine.
- Machines need to be retooled (sometimes) and reset by an operative.
- Operatives may be away attending other tasks.

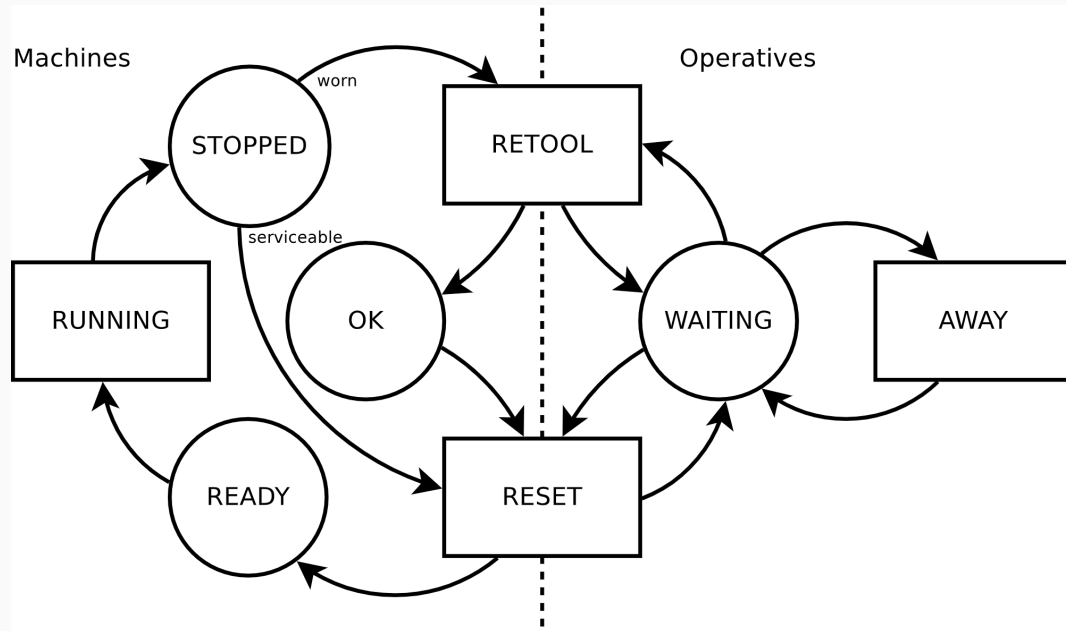


Example: Job Shop



From M. Pidd (1988), Section 5.3.1:

- Jobs are allocated to the first available machine.
- Machines need to be retooled (sometimes) and reset by an operative.
- Operatives may be away attending other tasks.




```
job ← trajectory() %>%
  seize("machine") %>%
  timeout(RUNNING) %>%
  branch(
    CHECK_WORN, continue = TRUE,
    trajectory() %>%
      seize("operative") %>%
      timeout(RETOOL) %>%
      release("operative")
  ) %>%
  seize("operative") %>%
  timeout(RESET) %>%
  release("operative") %>%
  release("machine")
```

Example: Job Shop



From M. Pidd (1988), Section 5.3.1:

- Jobs are allocated to the first available machine.
- Machines need to be retooled (sometimes) and reset by an operative.
- Operatives may be away attending other tasks.



```
job ← trajectory() %>%
  seize("machine") %>%
  timeout(RUNNING) %>%
  branch(
    CHECK_WORN, continue = TRUE,
    trajectory() %>%
      seize("operative") %>%
      timeout(RETOOL) %>%
      release("operative")
  ) %>%
  seize("operative") %>%
  timeout(RESET) %>%
  release("operative") %>%
  release("machine")
```

```
task ← trajectory() %>%
  seize("operative") %>%
  timeout(AWAY) %>%
  release("operative")
```

Example: Job Shop



```
library(simmer); set.seed(1234)

RUNNING ← function() rexp(1, 1)
CHECK_WORN ← function() runif(1) < 0.2
RETOOL ← function() rexp(1, 2)
RESET ← function() rexp(1, 3)
AWAY ← function() rexp(1, 1)
```

```
job ← trajectory() %>%
  seize("machine") %>%
  timeout(RUNNING) %>%
  branch(
    CHECK_WORN, continue = TRUE,
    trajectory() %>%
      seize("operative") %>%
      timeout(RETOOL) %>%
      release("operative")
  ) %>%
  seize("operative") %>%
  timeout(RESET) %>%
  release("operative") %>%
  release("machine")
```

```
task ← trajectory() %>%
  seize("operative") %>%
  timeout(AWAY) %>%
  release("operative")
```

Example: Job Shop



```
library(simmer); set.seed(1234)

RUNNING ← function() rexp(1, 1)
CHECK_WORN ← function() runif(1) < 0.2
RETOOL ← function() rexp(1, 2)
RESET ← function() rexp(1, 3)
AWAY ← function() rexp(1, 1)

job ← trajectory() %>%
  ...
task ← trajectory() %>%
  ...

NEW_JOB ← function() rexp(1, 5)
NEW_TASK ← function() rexp(1, 1)
```

Example: Job Shop



```
library(simmer); set.seed(1234)

RUNNING ← function() rexp(1, 1)
CHECK_WORN ← function() runif(1) < 0.2
RETOOL ← function() rexp(1, 2)
RESET ← function() rexp(1, 3)
AWAY ← function() rexp(1, 1)

job ← trajectory() %>%
  ...
task ← trajectory() %>%
  ...

NEW_JOB ← function() rexp(1, 5)
NEW_TASK ← function() rexp(1, 1)
```

```
env ← simmer("Job Shop") %>%
  add_resource("machine", 10) %>%
  add_resource("operative", 5) %>%
  add_generator("job", job, NEW_JOB) %>%
  add_generator("task", task, NEW_TASK) %>%
  run(until=1000)
env
```

```
## simmer environment: Job Shop | now: 1000 | next: 1000.118
## { Monitor: in memory }
## { Resource: machine | monitored: TRUE | server status: 90
## { Resource: operative | monitored: TRUE | server status:
## { Source: job | monitored: 1 | n_generated: 4954 }
## { Source: task | monitored: 1 | n_generated: 1041 }
```


Example: Job Shop



```
library(simmer); set.seed(1234)

RUNNING ← function() rexp(1, 1)
CHECK_WORN ← function() runif(1) < 0.2
RETOOL ← function() rexp(1, 2)
RESET ← function() rexp(1, 3)
AWAY ← function() rexp(1, 1)

job ← trajectory() %>%
  ...
task ← trajectory() %>%
  ...

NEW_JOB ← function() rexp(1, 5)
NEW_TASK ← function() rexp(1, 1)
```

```
env ← simmer("Job Shop") %>%
  add_resource("machine", 10) %>%
  add_resource("operative", 5) %>%
  add_generator("job", job, NEW_JOB) %>%
  add_generator("task", task, NEW_TASK) %>%
  run(until=1000)
env
```

```
## simmer environment: Job Shop | now: 1000 | next: 1000.118
## { Monitor: in memory }
## { Resource: machine | monitored: TRUE | server status: 90
## { Resource: operative | monitored: TRUE | server status:
## { Source: job | monitored: 1 | n_generated: 4954 }
## { Source: task | monitored: 1 | n_generated: 1041 }
```

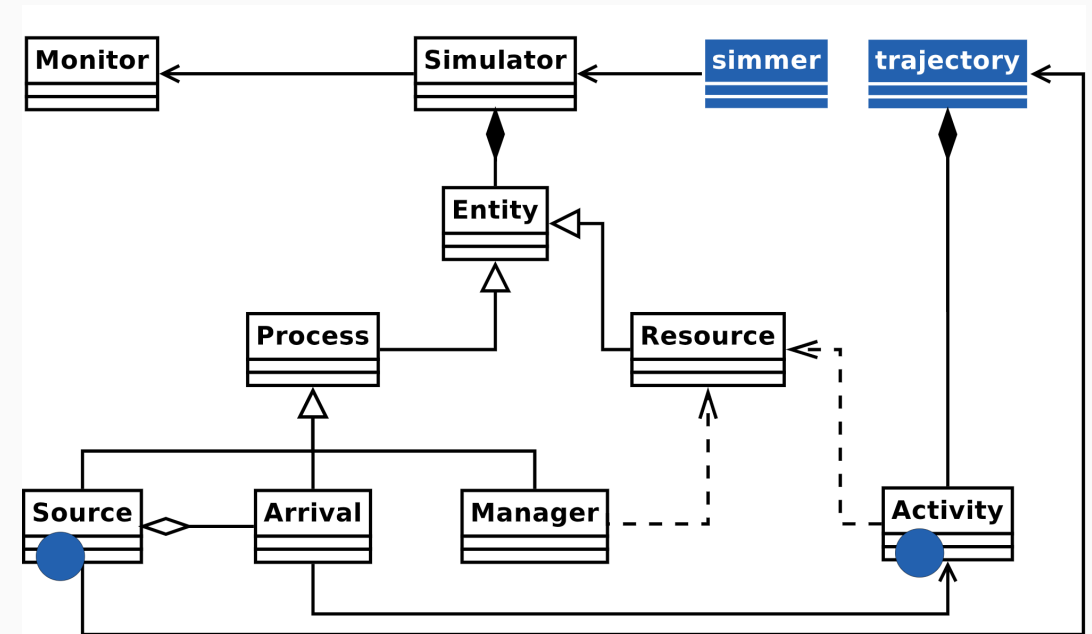
```
aggregate(cbind(server, queue) ~ resource, get_mon_resources(env), mean)
```

```
##   resource  server  queue
## 1  machine 7.632616 0.5598666
## 2 operative 3.389082 0.3436009
```



Terminology

Architecture



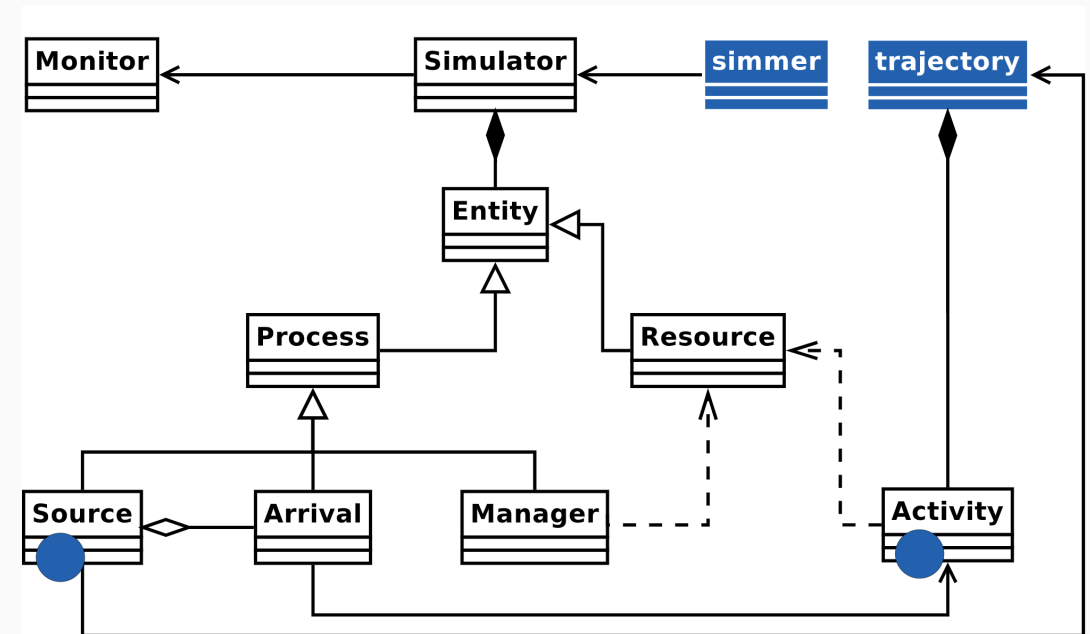
Overview of the C++ core (white) and R API (blue)



Terminology

- **Resource:** server (configurable capacity) + priority queue (configurable size), supports preemption
- **Manager:** modifies resources at run time (schedule)

Architecture



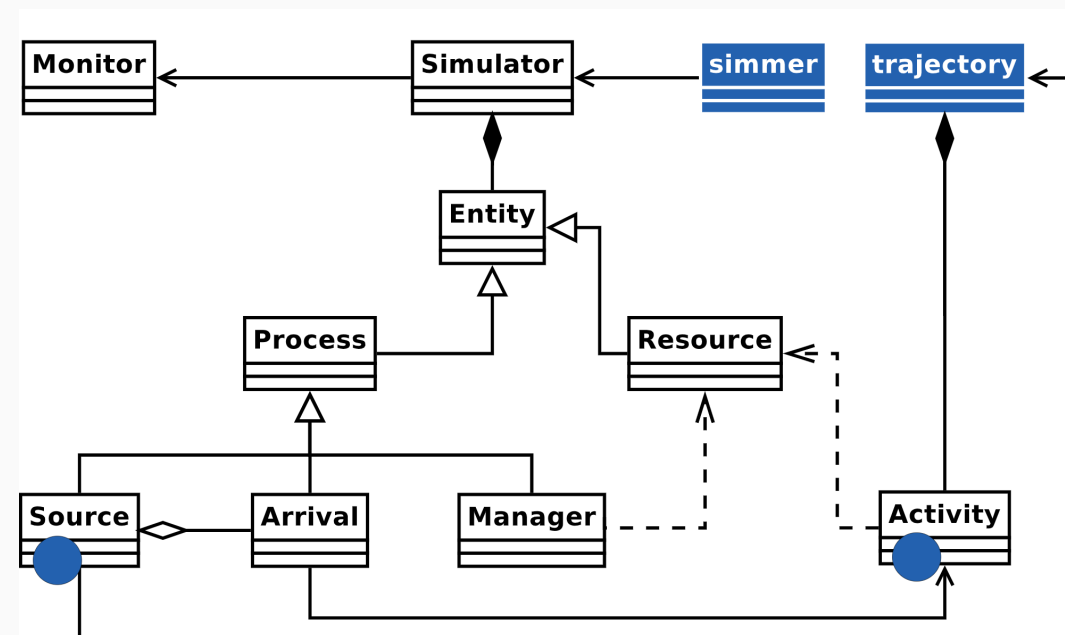
Overview of the C++ core (white) and R API (blue)



Terminology

- **Resource:** server (configurable capacity) + priority queue (configurable size), supports preemption
- **Manager:** modifies resources at run time (schedule)
- **Source:** creates new arrivals following some distribution of inter-arrival times
- **Arrival:** interacting processes, with attributes and prioritization values

Architecture



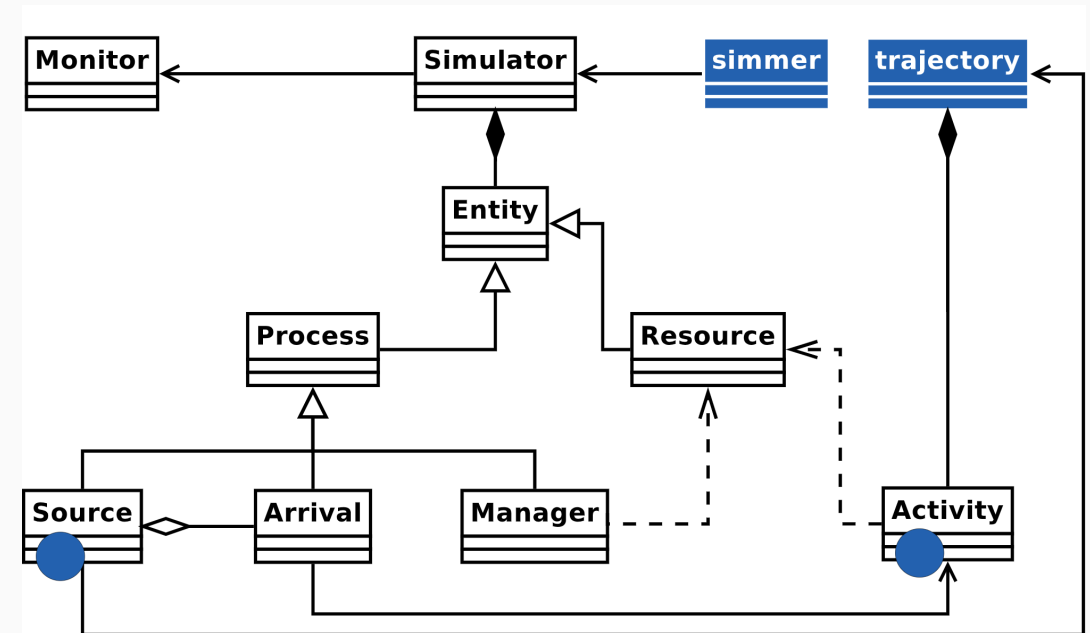
Overview of the C++ core (white) and R API (blue)



Terminology

- **Resource:** server (configurable capacity) + priority queue (configurable size), supports preemption
- **Manager:** modifies resources at run time (schedule)
- **Source:** creates new arrivals following some distribution of inter-arrival times
- **Arrival:** interacting processes, with attributes and prioritization values
- **Trajectory:** interlinkage of activities, a common path for arrivals of the same type
- **Activity:** unit of action, main building block

Architecture



Overview of the C++ core (white) and R API (blue)



Trajectory

Similar to `dplyr` for data manipulation. In the words of H. Wickham,

by constraining your options, it simplifies how you can think about [something]

- Trajectories are recipes, lists of activities defining the life time of arrivals
- Activities are common functional DES blocks



Trajectory

Similar to `dplyr` for data manipulation. In the words of H. Wickham,

by constraining your options, it simplifies how you can think about [something]

- Trajectories are recipes, lists of activities defining the life time of arrivals
- Activities are common functional DES blocks

Fixed vs. dynamic parameters:

```
traj0 ← trajectory() %>%  
  log_("Entering the trajectory") %>%  
  timeout(10) %>%  
  log_("Leaving the trajectory")
```

```
traj1 ← trajectory() %>%  
  log_(function() "Entering the trajectory") %>%  
  timeout(function() 10) %>%  
  log_(function() "Leaving the trajectory")
```



Trajectory

Activities

- Spend time in the system
 - `timeout`, `timeout_from_attribute`, `timeout_from_global`
- Modify arrival properties
 - `set_attribute`, `set_global`
 - `set_prioritization`
- Interaction with resources
 - `seize`, `release`
 - `set_capacity`, `set_queue_size`
 - `select`, `seize_selected...`
- Interaction with sources
 - `activate`, `deactivate`
 - `set_trajectory`, `set_source`
- Loops
 - `rollback`
- Branching
 - `branch`
 - `clone`, `synchronize`
- Batching
 - `batch`, `separate`
- Inter-arrival communication
 - `send`, `trap`, `untrap`, `wait`
- Reneging
 - `leave`
 - `renege_in`, `renege_if`, `renege_abort`
 - `handle_unfinished`
- Debugging
 - `log_`
 - `stop_if`



Trajectory

Activities

- Spend time in the system
 - `timeout`, `timeout_from_attribute`, `timeout_from_global`
- Modify arrival properties
 - `set_attribute`, `set_global`
 - `set_prioritization`
- Interaction with resources
 - `seize`, `release`
 - `set_capacity`, `set_queue_size`
 - `select`, `seize_selected...`
- Interaction with sources
 - `activate`, `deactivate`
 - `set_trajectory`, `set_source`
- Loops
 - `rollback`
- Branching
 - `branch`
 - `clone`, `synchronize`
- Batching
 - `batch`, `separate`
- Inter-arrival communication
 - `send`, `trap`, `untrap`, `wait`
- Reneging
 - `leave`
 - `renege_in`, `renege_if`, `renege_abort`
 - `handle_unfinished`
- Debugging
 - `log_`
 - `stop_if`

(Plus many getters to retrieve parameters at run time)



Trajectory

Activities

Simulation
environment

- Create a simulator and attach a monitor
 - `simmer`
 - `monitor_mem` (default), `monitor_csv`... (extensible, see `simmer.mon` on GitHub)



Trajectory

Activities

Simulation environment

- Create a simulator and attach a monitor
 - `simmer`
 - `monitor_mem` (default), `monitor_csv`... (extensible, see `simmer.mon` on GitHub)
- Add sources of arrivals
 - `add_generator`: based on a distribution function
 - `add_dataframe`: based on a data frame (additional columns as attributes)



Trajectory

Activities

Simulation environment

- Create a simulator and attach a monitor
 - `simmer`
 - `monitor_mem` (default), `monitor_csv` ... (extensible, see `simmer.mon` on GitHub)
- Add sources of arrivals
 - `add_generator`: based on a distribution function
 - `add_dataframe`: based on a data frame (additional columns as attributes)
- Add resources
 - `add_resource`: priority resource, with capacity and queue size; optional preemption



Trajectory

Activities

Simulation environment

- Create a simulator and attach a monitor
 - `simmer`
 - `monitor_mem` (default), `monitor_csv`... (extensible, see `simmer.mon` on GitHub)
- Add sources of arrivals
 - `add_generator`: based on a distribution function
 - `add_dataframe`: based on a data frame (additional columns as attributes)
- Add resources
 - `add_resource`: priority resource, with capacity and queue size; optional preemption
- Add global variables
 - `add_global`



Trajectory

Activities

Simulation environment

- Create a simulator and attach a monitor
 - `simmer`
 - `monitor_mem` (default), `monitor_csv`... (extensible, see `simmer.mon` on GitHub)
- Add sources of arrivals
 - `add_generator`: based on a distribution function
 - `add_dataframe`: based on a data frame (additional columns as attributes)
- Add resources
 - `add_resource`: priority resource, with capacity and queue size; optional preemption
- Add global variables
 - `add_global`
- Run the simulation
 - `run`, `stepn`



Trajectory

Activities

Simulation
environment

Monitoring

`simmer` automatically records every change in the state of the system. All these statistics can be retrieved after the simulation:

```
names( get_mon_arrivals(simmer(), per_resource=FALSE) )
```

```
## [1] "name"          "start_time"    "end_time"      "activity_time"  
## [5] "finished"
```

```
names( get_mon_arrivals(simmer(), per_resource=TRUE) )
```

```
## [1] "name"          "start_time"    "end_time"      "activity_time"  
## [5] "resource"
```

```
names( get_mon_attributes(simmer()) )
```

```
## [1] "time"  "name"  "key"   "value"
```

```
names( get_mon_resources(simmer()) )
```

```
## [1] "resource"  "time"      "server"     "queue"      "capacity"  
## [6] "queue_size" "system"    "limit"
```



Queueing systems

Natural way to simulate CTMC and birth-death processes:

```
set.seed(1234)
lambda ← 2
mu ← 4
rho ← lambda/mu
```

```
mm1.traj ← trajectory() %>%
  seize("mm1.resource", 1) %>%
  timeout(function() rexp(1, mu)) %>%
  release("mm1.resource", 1)
```

```
mm1.env ← simmer() %>%
  add_resource("mm1.resource", 1, Inf) %>%
  add_generator("arrival", mm1.traj,
               function() rexp(1, lambda)) %>%
  run(2000)
```



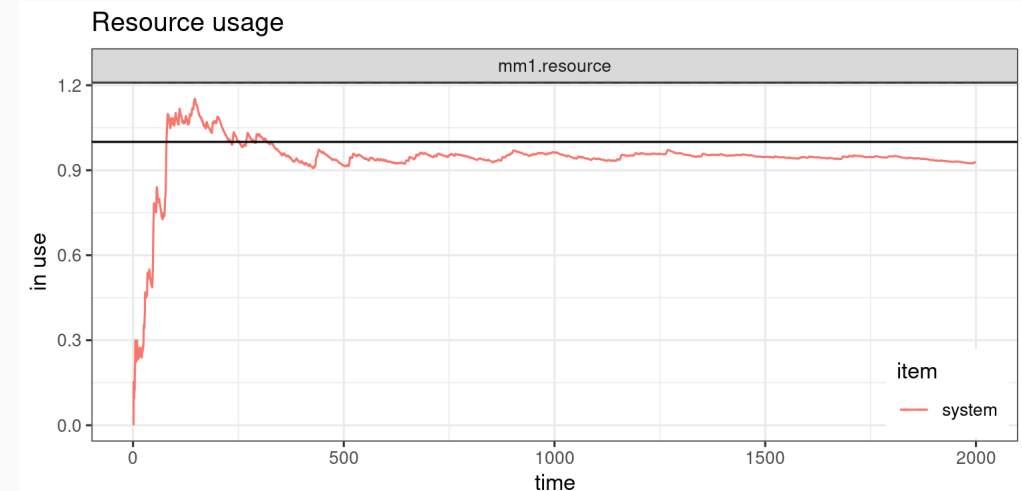
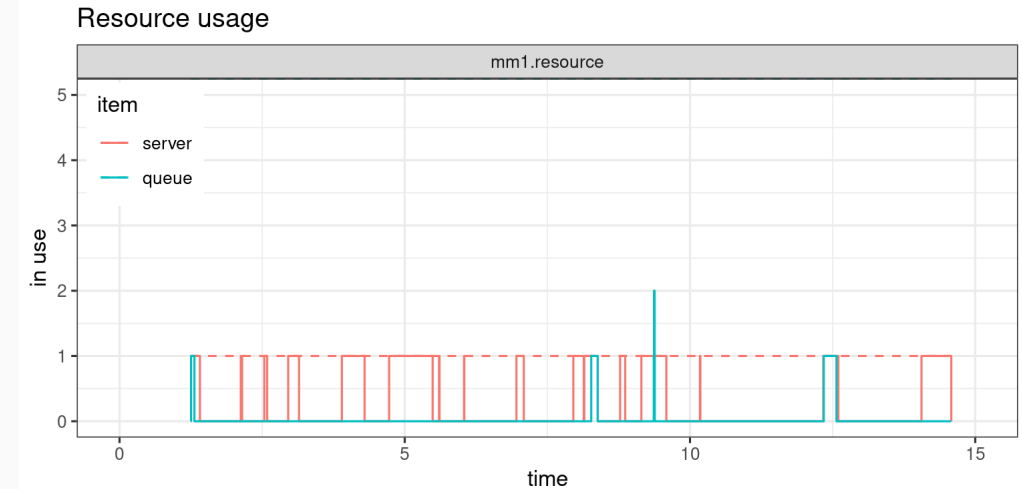

Queueing systems

Natural way to simulate CTMC and birth-death processes:

```
set.seed(1234)
lambda <- 2
mu <- 4
rho <- lambda/mu
```

```
mm1.traj <- trajectory() %>%
  seize("mm1.resource", 1) %>%
  timeout(function() rexp(1, mu)) %>%
  release("mm1.resource", 1)
```

```
mm1.env <- simmer() %>%
  add_resource("mm1.resource", 1, Inf) %>%
  add_generator("arrival", mm1.traj,
    function() rexp(1, lambda)) %>%
  run(2000)
```





Replication

Easy replication with standard R functions:

```
mm1.envs <- lapply(1:20, function(i) {  
  simmer() %>%  
    add_resource("mm1.resource", 1, Inf) %>%  
    add_generator("arrival", mm1.traj,  
                 function() rexp(100, lambda)) %>%  
    run(1000/lambda)  
})
```



Replication

Easy replication with standard R functions:

```
mm1.envs <- lapply(1:20, function(i) {  
  simmer() %>%  
    add_resource("mm1.resource", 1, Inf) %>%  
    add_generator("arrival", mm1.traj,  
                  function() rexp(100, lambda)) %>%  
    run(1000/lambda)  
})
```

Parallelization

Even easier parallelization of replicas:

```
mm1.envs <- parallel::mclapply(1:20, function(i) {  
  simmer() %>%  
    add_resource("mm1.resource", 1, Inf) %>%  
    add_generator("arrival", mm1.traj,  
                  function() rexp(100, lambda)) %>%  
    run(1000/lambda) %>%  
    wrap()  
}, mc.cores=4)
```



Replication

Easy replication with standard R functions:

```
mm1.envs <- lapply(1:20, function(i) {  
  simmer() %>%  
    add_resource("mm1.resource", 1, Inf) %>%  
    add_generator("arrival", mm1.traj,  
                 function() rexp(100, lambda)) %>%  
    run(1000/lambda)  
})
```

```
head(get_mon_arrivals(mm1.envs), 3)
```

##	name	start_time	end_time	activity_time	finished	replication
## 1	arrival0	0.1455818	0.4810091	0.3354273	TRUE	1
## 2	arrival1	0.2380988	0.6834863	0.2024772	TRUE	1
## 3	arrival2	0.4526185	0.8724942	0.1890079	TRUE	1

Parallelization

Even easier parallelization of replicas:

```
mm1.envs <- parallel::mclapply(1:20, function(i) {  
  simmer() %>%  
    add_resource("mm1.resource", 1, Inf) %>%  
    add_generator("arrival", mm1.traj,  
                 function() rexp(100, lambda)) %>%  
    run(1000/lambda) %>%  
    wrap()  
}, mc.cores=4)
```



Best practices

There are usually multiple valid ways of mapping the identified resources and processes into the `simmer` API



Best practices

There are usually multiple valid ways of mapping the identified resources and processes into the `simmer` API

Design pattern 1

```
beep ← trajectory() %>%  
  log_("beeeep!")
```

```
env ← simmer() %>%  
  add_generator("beep", beep, function() 1) %>%  
  run(2.5)
```

```
## 1: beep0: beeeep!
```

```
## 2: beep1: beeeep!
```



Best practices

There are usually multiple valid ways of mapping the identified resources and processes into the `simmer` API

Design pattern 1

```
beep ← trajectory() %>%  
  log_("beeeep!")
```

```
env ← simmer() %>%  
  add_generator("beep", beep, function() 1) %>%  
  run(2.5)
```

```
## 1: beep0: beeeep!
```

```
## 2: beep1: beeeep!
```

Design pattern 2

```
alarm ← trajectory() %>%  
  timeout(1) %>%  
  log_("beeeep!") %>%  
  rollback(2)
```

```
env ← simmer() %>%  
  add_generator("alarm", alarm, at(0)) %>%  
  run(2.5)
```

```
## 1: alarm0: beeeep!
```

```
## 2: alarm0: beeeep!
```



Comparison with similar frameworks (out-of-date!):

- SimPy 3.0.9, Python 2.7
- SimJulia 0.3.14, Julia 0.5.1

Heavy M/M/1, $\rho \approx 0.9$:

```
test_mm1_simmer ← function(n, m, mon=FALSE) {  
  mm1 ← trajectory() %>%  
    seize("server", 1) %>%  
    timeout(function() rexp(1, 1.1)) %>%  
    release("server", 1)  
  
  env ← simmer() %>%  
    add_resource("server", 1, mon=mon) %>%  
    add_generator("customer", mm1,  
                 function() rexp(m, 1), mon=mon) %>%  
    run(until=n)  
}
```


Performance





Comparison with similar frameworks (out-of-date!):

- SimPy 3.0.9, Python 2.7
- SimJulia 0.3.14, Julia 0.5.1

Heavy M/M/1, $\rho \approx 0.9$:

```
test_mm1_simmer <- function(n, m, mon=FALSE) {  
  mm1 <- trajectory() %>%  
    seize("server", 1) %>%  
    timeout(function() rexp(1, 1.1)) %>%  
    release("server", 1)  
  
  env <- simmer() %>%  
    add_resource("server", 1, mon=mon) %>%  
    add_generator("customer", mm1,  
                 function() rexp(m, 1), mon=mon) %>%  
    run(until=n)  
}
```





The cost of calling R from C++ revisited

Very simple deterministic test to study the impact:

```
test_simmer <- function(n, delay) {  
  test <- trajectory() %>%  
    timeout(delay)  
  simmer() %>%  
    add_generator("test", test, at(1:n)) %>%  
    run() %>%  
    get_mon_arrivals()  
}  
  
test_simmer(5, 1)[,1:5]
```

##	name	start_time	end_time	activity_time	finished
## 1	test0	1	2	1	TRUE
## 2	test1	2	3	1	TRUE
## 3	test2	3	4	1	TRUE
## 4	test3	4	5	1	TRUE
## 5	test4	5	6	1	TRUE



The cost of calling R from C++ revisited

Very simple deterministic test to study the impact:

```
test_simmer ← function(n, delay) {  
  test ← trajectory() %>%  
    timeout(delay)  
  simmer() %>%  
    add_generator("test", test, at(1:n)) %>%  
    run() %>%  
    get_mon_arrivals()  
}  
  
test_simmer(5, 1)[,1:5]
```

```
##      name start_time end_time activity_time finished  
## 1 test0           1         2             1      TRUE  
## 2 test1           2         3             1      TRUE  
## 3 test2           3         4             1      TRUE  
## 4 test3           4         5             1      TRUE  
## 5 test4           5         6             1      TRUE
```

Original benchmark in the JSS paper:

Expr	Min	Mean	Median	Max
test_simmer(n, 1)	429.8663	492.365	480.5408	599.3547
test_simmer(n, function() 1)	3067.9957	3176.963	3165.6859	3434.7979
test_R_for(n)	2053.0840	2176.164	2102.5848	2438.6836



The cost of calling R from C++ revisited

Very simple deterministic test to study the impact:

```
test_simmer <- function(n, delay) {  
  test <- trajectory() %>%  
    timeout(delay)  
  simmer() %>%  
    add_generator("test", test, at(1:n)) %>%  
    run() %>%  
    get_mon_arrivals()  
}  
  
test_simmer(5, 1)[,1:5]
```

```
##      name start_time end_time activity_time finished  
## 1 test0           1         2             1      TRUE  
## 2 test1           2         3             1      TRUE  
## 3 test2           3         4             1      TRUE  
## 4 test3           4         5             1      TRUE  
## 5 test4           5         6             1      TRUE
```

Original benchmark in the JSS paper:

Expr	Min	Mean	Median	Max
test_simmer(n, 1)	429.8663	492.365	480.5408	599.3547
test_simmer(n, function() 1)	3067.9957	3176.963	3165.6859	3434.7979
test_R_for(n)	2053.0840	2176.164	2102.5848	2438.6836

Update with `-DRCPP_USE_UNWIND_PROTECT`:

Expr	Min	Mean	Median	Max
test_simmer(n, 1)	467.8971	481.213	476.1667	521.4916
test_simmer(n, function() 1)	498.2631	583.777	561.6798	816.1343
test_R_for(n)	1158.9348	1201.460	1196.7223	1244.4041



- Generic yet powerful process-oriented Discrete-Event Simulation framework for R [1, 2]
- Combines a robust and **fast** C++ simulation core with a **rich and flexible** R API

[1] **Ucar I**, Smeets B, Azcorra A (2019). “simmer: Discrete-Event Simulation for R.” *Journal of Statistical Software*, 90(2), 1-30. doi: [10.18637/jss.v090.i02](https://doi.org/10.18637/jss.v090.i02).

[2] **Ucar I**, Hernández JA, Serrano P, Azcorra A (2018). “Design and Analysis of 5G Scenarios with simmer: An R Package for Fast DES Prototyping.” *IEEE Communications Magazine*, 56(11), 145-151. doi: [10.1109/MCOM.2018.1700960](https://doi.org/10.1109/MCOM.2018.1700960).



- Generic yet powerful process-oriented Discrete-Event Simulation framework for R [1, 2]
- Combines a robust and **fast** C++ simulation core with a **rich and flexible** R API
- Broad set of **activities**, the basic building block; extensible via custom routines
- Activities are chained into a **trajectory**, a common path for processes of the same type

[1] **Ucar I**, Smeets B, Azcorra A (2019). “simmer: Discrete-Event Simulation for R.” *Journal of Statistical Software*, 90(2), 1-30. doi: [10.18637/jss.v090.i02](https://doi.org/10.18637/jss.v090.i02).

[2] **Ucar I**, Hernández JA, Serrano P, Azcorra A (2018). “Design and Analysis of 5G Scenarios with simmer: An R Package for Fast DES Prototyping.” *IEEE Communications Magazine*, 56(11), 145-151. doi: [10.1109/MCOM.2018.1700960](https://doi.org/10.1109/MCOM.2018.1700960).



- Generic yet powerful process-oriented Discrete-Event Simulation framework for R [1, 2]
- Combines a robust and **fast** C++ simulation core with a **rich and flexible** R API
- Broad set of **activities**, the basic building block; extensible via custom routines
- Activities are chained into a **trajectory**, a common path for processes of the same type
- **Automatic monitoring**: focus on modelling

[1] **Ucar I**, Smeets B, Azcorra A (2019). “simmer: Discrete-Event Simulation for R.” *Journal of Statistical Software*, 90(2), 1-30. doi: [10.18637/jss.v090.i02](https://doi.org/10.18637/jss.v090.i02).

[2] **Ucar I**, Hernández JA, Serrano P, Azcorra A (2018). “Design and Analysis of 5G Scenarios with simmer: An R Package for Fast DES Prototyping.” *IEEE Communications Magazine*, 56(11), 145-151. doi: [10.1109/MCOM.2018.1700960](https://doi.org/10.1109/MCOM.2018.1700960).



- Generic yet powerful process-oriented Discrete-Event Simulation framework for R [1, 2]
- Combines a robust and **fast** C++ simulation core with a **rich and flexible** R API
- Broad set of **activities**, the basic building block; extensible via custom routines
- Activities are chained into a **trajectory**, a common path for processes of the same type
- **Automatic monitoring**: focus on modelling
- **Integration**: easy replication, parallelization, analysis...

[1] **Ucar I**, Smeets B, Azcorra A (2019). “simmer: Discrete-Event Simulation for R.” *Journal of Statistical Software*, 90(2), 1-30. doi: [10.18637/jss.v090.i02](https://doi.org/10.18637/jss.v090.i02).

[2] **Ucar I**, Hernández JA, Serrano P, Azcorra A (2018). “Design and Analysis of 5G Scenarios with simmer: An R Package for Fast DES Prototyping.” *IEEE Communications Magazine*, 56(11), 145-151. doi: [10.1109/MCOM.2018.1700960](https://doi.org/10.1109/MCOM.2018.1700960).

Thanks, and happy **simmering**!

<https://r-simmer.org>

